

Module 3: Feedforward Neural Networks

Lesson 1: Feedforward Neural Networks

単語

- congruent: \cong , if one thing is congruent with another thing, they are similar or fit together well.

内容

- The mode of action of feedforward neural networks.
- The mathematical equations of the hidden layers, a building block that **makes neural networks special** among machine learning models.

Feedforward Neural Networks

- A Feedforward Neural Network defines a mapping from input x to output y as: $y = f(x; \theta)$.
 - θ : **learned** parameters.
 - The concept of learned parameters is important, as we do not start with the correct form of the function f , which maps the inputs to outputs directly.
 - This means that neural networks can be thought of as **function approximators**.
- An N layer FNN is represented as the function composition:

$$f(x; \theta) = f^{(N)} \left(f^{(N-1)} \left(\dots f^{(2)} \left(f^{(1)}(x) \right) \right) \right).$$

- N is a large number.
- This layering led to the name deep learning for the field describing these sequences of functions.
- Feedforward: This is because information flows from the input x through some intermediate steps, all the way to the output y **without any feedback connections**.
 - The terms are used in the same way when describing **control** for self-driving cars.
- Functions to estimate
 - Object Classification: Image \rightarrow Label.
 - Object Detection: Image \rightarrow Label + Location.
 - **Depth Estimation**: Image \rightarrow Depth for every pixel.
 - Semantic Segmentation: Image \rightarrow **Label for every pixel**.
- This flexibility to represent **hard-to-model processes** is what makes neural networks so popular.

Mode of Action of Neural Networks

- Training: Given neural network examples of $f^*(x)$ for a **wide variation** of the input x .
 - Then optimize its parameters θ to force $f(x; \theta) \cong f^*(x)$.
- Pairs of x and $f^*(x)$ are called training data.
- Only output is specified by training data!
 - Network is free to do anything with its hidden layers.

Hidden Units

$$h_n = g(W^T h_{n-1} + b)$$

- Activation function g .
 - element-wise non-linear function.
- Parameters θ are the weights and biases of all the layers of the network.
- **Most of the time, g does not contain parameters to be learned** by the network.

The Rectified Linear Unit: ReLU

- The ReLU hidden unit is currently the default choice of activation function for Feedforward Neural Networks: $g(z) = \max(0, z)$.
- Since they are very similar to linear units, they're quite easy to optimize.

他のActivation Functions

- Sigmoid non-linearity.
- the hyperbolic tangent non-linearity. (Tanh)
- the generalization of ReLU: the maxout non-linearity.

Lesson 2: Output Layers and Loss Functions

内容

- The general process of designing machine learning algorithm, and extend it to the design of neural networks.
- Different types of neural networks loss functions that can be used depending on the type of task at hand.

Machine Learning Algorithm Design

- Generally, supervised machine learning models including neural networks have two **modes of operation, inference and training**.
 - Inference is usually the operation we used when we **deploy** the machine learning algorithms in the real world.
- For self-driving, the training data often takes the form of human annotated images which take a long time to produce.
- The Optimization procedure takes in the output of the loss function and provides a new set of parameters θ that provide a lower value for that loss function.
- The major difference between the design of traditional machine learning algorithms and the design of artificial neural networks, is that the **neural network only interacts with the loss function via the output layer**.
 - Therefore, it is quite reasonable that the **output layer and the loss function are designed together** depending on the task at hand.

Tasks: Classification and Regression

- Classification: Given input x , map it to one of k classes or categories.
 - Image classification, semantic segmentation.
- Regression: Given input x , **map it to a real number**.
 - Depth prediction, bounding box estimation.

Classification: Softmax Output Layers (大事)

- Softmax output layers are most often used as the output of a classifier, to represent the probability distribution over K different classes.
- The Softmax output layer is comprised of:
 - A linear transformation: $z = W^T h + b$.
 - Followed by the Softmax function: $Softmax(z_i) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$.

Classification: Cross-Entropy Loss Function (大事)

- By **considering** the output of the softmax output layer as a probability distribution, the Cross Entropy Loss function is derived using maximum likelihood as:

$$L(\theta) = -\log(Softmax(z_i)) = -z_i + \log \sum_j \exp(z_j).$$
- The Cross-Entropy Loss has two terms to control how close the output of the network is to the true probability.
- z_i is usually called the **class logit**, which comes from the field of logistic regression.

- When minimizing this loss function, the negative of the class logit z_i (つまり $-z_i$) encourages the network to output a large value for the probability of the correct class.
- The second term on the other hand, encourages the output of the affine transformation (つまり z_i) to be small.
- Note how the loss function heavily penalizes erroneous predictions even when the difference in output is only one.
 - これは $-z_i + \log \sum_j \exp(z_j)$ からすぐ分かれるでしょう。分母はlog、分子はlinearだから。正しさも間違いも全部exaggerateされてる。
- This difference accelerates the learning process and rapidly steers network outputs to the true values during training.

Regression: Linear Output Layers

- Linear Output Units are based only on an affine transformation with no non-linearity: $z = W^T h + b$.
- Linear Output Units are usually used with the Mean Squared Error loss function to model the mean of a probability distribution: $L(\theta) = \sum_i (z_i - f^*(x_i))^2$.

Lesson 3: Neural Network Training with Gradient Descent

内容

- How to train a neural network using the iterative optimization algorithm: gradient descent.
- How to initialize parameters at the start of the optimization process.

Neural Network Loss Functions

- Thousands of training example pairs $[x, f^*(x)]$.
- The Loss function computed over all N training examples is termed the Training Loss and can be written as: $J(\theta) = \frac{1}{N} \sum_{i=1}^N L[f(x_i, \theta), f^*(x_i)]$.

- The gradient of the training loss with respect to the parameters θ can be written as:

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \left[\frac{1}{N} \sum_{i=1}^N L[f(x_i, \theta), f^*(x_i)] \right] = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)].$$

- using the fact that the gradient and the sum are linear operators.

Batch Gradient Descent

- Batch Gradient Descent is an iterative first order optimization procedure.
- Batch Gradient Descent Algorithm:

- Initialize parameters θ .
- While Stopping Condition is Not Met
 - Compute gradient of loss function over all training examples:

$$\nabla_{\theta} J(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)].$$

- Update parameters according to: $\theta \leftarrow \theta - \epsilon \nabla_{\theta} J(\theta)$.

- How do we initialize the parameters θ and how do we decide when to actually stop the algorithm?

- The answer to both of these questions is still highly based on heuristics that work well in practice.

Parameter Initialization and Stopping Conditions (大事)

- Parameter Initialization
 - Weights: initialized by randomly sampling from a standard normal distribution.
 - Biases: initialized to 0.
 - Other heuristics exist.
- Stopping Conditions
 - Number of iterations.
 - Change in θ values: Stop if $\theta_{new} - \theta_{old} < \text{Threshold}$.
 - Change in $J(\theta)$ value: Stop if $J(\theta_{new}) - J(\theta_{old}) < \text{Threshold}$.
 - J is N training samples of L of the average value.

Batch Gradient Descent

- Backpropagation used to compute $\nabla_{\theta} J(\theta)$ is very expensive to compute over the whole training dataset.
 - 原因: Backpropagation involves computing the output of the network for the example on which we would like to evaluate the gradient.
- $\nabla_{\theta} J(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)]$ is a mean!
- Standard error of the mean estimated from N samples is $\frac{\sigma}{\sqrt{N}}$, where σ is the standard deviation of the value of the samples.
- Using all samples to estimate the gradient results is less than linear return in accuracy of this estimate.
- Use a small subsample (Minibatch) of the training data to estimate the gradient!

Stochastic (Minibatch) Gradient Descent

- Stochastic (Minibatch) Gradient Descent Algorithm
 - Initialize parameters θ .
 - While Stopping Condition is False
 - Sample a preset number N' of examples (minibatch) from the training data.
 - new parameter N' .
 - Compute gradient of Loss Function over these N' training examples:
$$\nabla_{\theta} J(\theta) = \frac{1}{N'} \sum_{i=1}^{N'} \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)].$$
 - Update parameters according to: $\theta \leftarrow \theta - \epsilon \nabla_{\theta} J(\theta)$.

What Minibatch Size to Use? これは覚えてる

「Deep Learning」 p270~275: 8.1.3 Batch and Minibatch Algorithmsに全部書いてある!

- GPUs work better with powers of 2 batch sizes.
- Large batch sizes > 256:
 - Hardware underutilized with very small batch sizes.
 - More accurate estimate of the gradient, but with less than linear returns.
- Small batch size < 64:
 - Small batches can offer a regularizing effect. The best generalization error is often achieved with batch size of 1.
 - 「Deep Learning」 p272によると、Small batches can offer a regularizing effect, perhaps due to the noise they add to the learning process. Generalization error is often best for a batch size of 1. Training with such a small batch size might require a small learning rate to maintain stability because of the high variance in the estimate of the gradient. The total runtime can be very high as a result of the need to make more steps, both because of the reduced learning rate and because it takes more steps to observe the entire training set.

- Small batch sizes allow for faster convergence, as the algorithm can compute the parameter updates rapidly.
- Always make sure dataset is shuffled before sampling minibatch.
- As a result of these trade-offs, **typical power of two minibatch sizes range from 32 to 256**, with smaller sizes sometimes being attempted for large models or to improve generalization.

SGD Variations

- Many variations of SGD exist
 - Momentum SGD, Nesterov Momentum SGD.
 - Ada-Grad, RMS-Prop.
 - ADAM (Adaptive Moment Estimation).
- Which one to use?
 - **ADAM**: Implemented in most deep neural network libraries, **fairly robust to the choice of the learning rate and other hyperparameters**.

Lesson 4: Data Splits and Neural Network Performance Evaluation

内容

- How to split a dataset for an unbiased estimate of performance. 意味はまだ分かっていない。そうだ！ training set, validation set, test setの話でしょう！
 - unbiased estimate,つまりtest setがtraining setとindependent, identically distributedという意味でしょう。
- How to improve the performance of neural network by observing the difference in performance on the various data splits.

Data Splits

- Given a large enough neural network, we are almost guaranteed to get a very low training loss.
 - This is due to the **very large number of parameters** in a typical deep neural network allowing it to **memorize the training data** to a large extent given a large enough number of training iterations.
- Training Split: used to minimize the Loss Function.
- **Validation Split**: used to **choose best hyperparameters**, such as the learning rate, **number of layers, the number of units per layer**, and the activation function type.
 - Hyperparameters are those parameters that either **modify the network structure** or affect the training process. つまりhyperparametersこそネットワークを決めるでしょう！
- Test Split: the neural network never observes this set. The developer never uses this set in the design process.
 - used to get an unbiased estimate of performance.
- データのサイズによって各splitのpercentageは変わる。 (大事)
 - **~10,000**
 - Training, 60%
 - Validation, 20%
 - Testing, 20%
 - **~1,000,000**
 - Having 20% of the data in the validation and test sets is **unnecessary** as the validation and test would contain **far more samples than are needed for the purpose**.
 - Training, 98%
 - Validation, 1%
 - Testing, 1%

Behavior of Split Specific Loss Functions

	Training	Validation	Testing	
	$J(\theta)_{train}$	$J(\theta)_{val}$	$J(\theta)_{test}$	$J(\theta)_{Minimum}$

	Training	Validation	Testing	
Good Estimator	0.21	0.25	0.30	0.18
Underfitting	1.9	1.9	2.1	
Overfitting	0.21	2.05	2.1	

- Good Estimator: the **loss on the three sets** are **fairly consistent** and the loss is close to the minimum achievable loss on the entire task.
- Underfitting: the neural network fails to bring the training loss down.
- Overfitting: caused by the neural network optimizing its parameters to precisely **reproduce** the training data output.
- The gap between **training and validation loss** is called the **generalization gap**.

Reducing the Effect of Underfitting/Overfitting

- Underfitting: (Training loss is high)
 - Train longer
 - If the **architecture** is **suitable for the task** at hand.
 - More layers or more parameters per layer.
 - Change architecture.
 - Overfitting: (Generalization gap is large) (大事)
 - More training data.
 - Unfortunately, for self-driving cars, collecting training data is very expensive as it requires engineering time for data collection and **a tremendous amount of annotator time to properly define the true output**.
 - Regularization.
 - Change architecture.
- A dataset should be split to a training, a validation and a test split.
 - Observing the performance on each of these splits helps in determining why a neural network is not performing well in the real world.
 - Underfitting: Train longer or use a larger neural network.
 - Overfitting: Regularization.
 - A much more commonly faced scenario in self-driving car perception is overfitting.

Lesson 5: Neural Network Regularization

内容

- Remedy overfitting through various regularization strategies including:
 - Parameter norm penalties.
 - Dropout.
 - Early Stopping.

Toy Example

- New Design
 - 6 layer NN, 6 Hidden Units/Layer.
 - Train set Loss: 0.1.
 - Val set Loss: 0.45.
 - Minimum Loss achievable: 0.1.
- Overfitting is caused by the network learning the noise in the training data.
 - Because the neural network has so many parameters, it is able to curve out small regions in the space that correspond to the noisy training examples.

Parameter Norm Penalties

$$J(\theta)_{reg} = J(\theta) + \alpha \Omega(\theta)$$

- α is a hyperparameter that weights the relative contribution of the norm penalty to the value of the loss function.
- $\Omega(\theta)$ is a measure of **how large θ 's value is**, usually an L_p Norm.

- Usually only constrain the size of weights and not biases: $J(\theta)_{reg} = J(\theta) + \alpha \Omega(W)$. 覚えてる!
- This is motivated by the fact that the number of weights is much larger than the number of biases in the neural network.
- しかしDeep Learning 7.1 Parameter Norm Penalties page 223によると、理由はこれじゃない。理由は2つある
 - First, “Each weight specifies how two variables interact. Fitting the weight well requires observing both variables in a variety of conditions. Each bias controls only a single variable. This means that we **do not induce too much variance by leaving the biases unregularized.**”
 - Second, “**Regularizing the bias parameters can introduce a significant amount of underfitting.**”

L2-Norm Parameter Penalty

$$\Omega(W) = \frac{1}{2} W^T W = \frac{1}{2} \|W\|_2^2.$$

- Toy exampleの結果
 - train set loss: 0.1 -> 0.176.
 - val set loss: 0.45 -> 0.182.
 - In this case, the decrease in the generalization gap is higher than the increase in training set loss.

Dropout

1. $P_{keep} = 0.5$ (例えば) .
 - At every training iteration, this probability is used to choose a subset of the network nodes to keep in the network.
2. Evaluate the output y after **cutting** all the **connections coming out of these units**.
3. Multiply the final weights by P_{keep} at the end of training.
 - Dropout can be intuitively explained as forcing the model to learn with missing input and hidden units.
 - Or in other words, with different version of itself.
 - Computationally inexpensive but powerful regularization method.
 - Does not significantly limit the type of model or training procedure that can be used.
 - Dropout layers are practically implemented in all neural network libraries!
 - We recommend using dropout whenever you have **dense feedforward neural network layers**.

Early Stopping

- Early stopping ends training when the validation loss **keeps increasing** for a preset number of iterations or epochs.
- After stopping the training algorithm, the set of **parameters with the lowest validation loss** is returned. stopした時のparametersではない。
- Early stopping should not be used as a first choice for regularization.
 - As it also limits the training time, which may interfere with the overall network performance.

Lesson 6: Convolutional Neural Networks

内容

- How a neural network can use cross-correlation in its hidden layers instead of general matrix multiplication, to form ConvNets.
 - to tailor neural networks for **image** input data.
- The advantages of using ConvNets over traditional neural networks for processing images.


ConvNets

- Used for processing data defined on **grid**.
- 1D time series data (sampled at regular intervals), 2D images, 3D videos.
- Two major type of layers

- Convolution Layers.
- Pooling Layers.
- Example: VGG 16.

Fully Connected VS Convolutional Layers

- Fully connected layers connect each node output to every node input in the next layer.
 - This is **implemented in software** through **dense matrix multiplication**.
 - $h_n = g(W^T h_{n-1} + b)$.
- Although counter-intuitive, convolutional layers use **cross-correlation not convolutions** for their linear operator instead of general matrix multiplication. (大事)
 - $h_n = g(W * h_{n-1} + b)$.
 - The logic behind using cross-correlation is that **if the parameters are learned, it does not matter if we flip the output or not**.
 - Since we are learning the weights of the convolutional layer, the flipping does not affect our results at all. よく分かっていない。
 - This results in **sparse connectivity**.
 - Each input element to the convolutional layer only affects a few output elements, thanks to the use of a **limited size kernel** for the convolutional operation.



Cross-Correlation (大事)

- Width: horizontal dimension of input volume.
- Height: vertical dimension of input volume.
- Depth: number of channels of input volume.
- Padding size: essential to retain shape.
- Each filter is comprised of a set of weights and a single bias.
 - The number of channels of the kernel needs to correspond to the number of channels of the input volume.
 - In this case, we have three weight channels per filter corresponding to red, green, and blue channels of the input image.
 - input imageの3つchannelsとfilterの3つchannelsとfilterのbiasと一緒に1つvalueを計算する。
- つまりinput imageのchannel数とoutputの形は関係ない。
- Usually we have multiple filters per convolutional layer.
- Notice that we get one output channel per filter.
- Stride: The number of pixels the filter is moved in the vertical and horizontal direction.

Output Volume Shape

- Filters are size $m \times m$.
- Number of filters = K .
- Stride = S , Padding = P .
- $W_{out} = \frac{W_{in} - m + 2 \times P}{S} + 1$.
- $H_{out} = \frac{H_{in} - m + 2 \times P}{S} + 1$.
- $D_{out} = K$.
- When designing ConvNets, it is very important to know what size output layers you'll end up with.
 - As an example, you don't want to reduce the size of your output volume too much if you are trying to detect small traffic signs and traffic lights on road scenes.

Pooling Layers: Max Pooling

- Pooling helps make the representations become invariant to small translations of the input.

Output Volume Shape

- Pool size $n \times n$.
- Stride = S .
- $W_{out} = \frac{W_{in} - n}{S} + 1$.
- $H_{out} = \frac{H_{in} - n}{S} + 1$.
- $D_{out} = D_{in}$.

Advantages of ConvNets

- Convolutional neural networks are by design, a natural choice to process images.
- Convolutional layers have **less parameters** than fully connected layers, **reducing** the chances of **overfitting**.
 - through **parameters sharing** and allows ConvNets to operate on larger images.
 - パラメータが多すぎると、training dataを覚えちゃう。
- Convolutional layers use the same parameters to process every block of the image. **Along with** pooling layers, this leads to **translation invariance** which is particularly important for image understanding.
 - つまりtranslation invarianceの結果はpoolingからだけではなく、cross-correlationからも。
- ConvNets were one of the first neural network models to perform well at a time where other feedforward architectures failed.
 - In many ways, **ConvNets carry the torch for the rest of deep learning**, and **pave the way** to the **relatively new acceptance** of neural networks in general.
- ConvNets were one of the first neural network models to solve important commercial applications, such as handwritten digit recognition in the early 1990s. [LeCun et. al.]

学んだこと

1. (https://en.wikipedia.org/wiki/Deep_learning) によると
 - For a feedforward neural network, the depth of the CAPs is that of the network and is the number of hidden layers plus one (as the output layer is also parameterized).
 - For recurrent neural networks, in which a signal may propagate through a layer more than once, the CAP depth is potentially unlimited.
 - CAP: credit assignment path, the chain of transformations from input to output.
2. np.maximum
 - **element-wise** maximum of array elements.

- Compare two arrays and returns a new array containing the element-wise maxima. If one of the elements being compared is a NaN, then that element is returned. (つまり常にNaNをreturn)
If both elements are NaNs then the first is returned.
- The maximum is equivalent to `np.where(x1>=x2, x1, x2)` when neither `x1` nor `x2` are nans, but it is **faster** and does **proper broadcasting**.
- Broadcastingについて、このページ (<https://docs.scipy.org/doc/numpy/user/basics.broadcasting.html>) まだ細かく読んでいないが。また、このページ (<https://numpy.org/devdocs/user/theory.broadcasting.html>) も: Array Broadcasting in Numpy. (大事)